

Exp.No.: 3 Map Reduce program to process a weather dataset**AIM:**

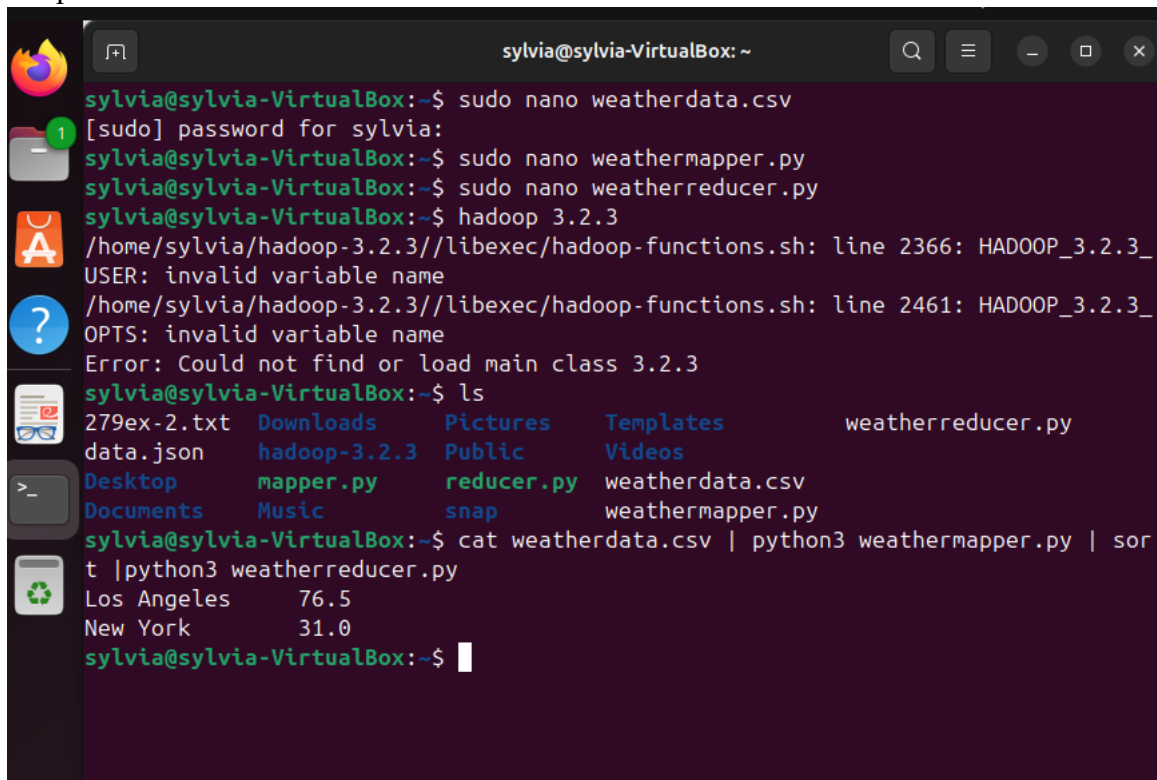
To implement MapReduce program to process a weather dataset.

Procedure:**Step 1: Create Data File:**

Create a file named "word_count_data.txt" and populate it with text data that you wish to analyse. Login with your hadoop user.

Download the dataset (weather data)

Output:



```
sylvia@sylvia-VirtualBox: ~
sylvia@sylvia-VirtualBox:~$ sudo nano weatherdata.csv
[sudo] password for sylvia:
sylvia@sylvia-VirtualBox:~$ sudo nano weathermapper.py
sylvia@sylvia-VirtualBox:~$ sudo nano weatherreducer.py
sylvia@sylvia-VirtualBox:~$ hadoop 3.2.3
/home/sylvia/hadoop-3.2.3//libexec/hadoop-functions.sh: line 2366: HADOOP_3.2.3_
USER: invalid variable name
/home/sylvia/hadoop-3.2.3//libexec/hadoop-functions.sh: line 2461: HADOOP_3.2.3_
OPTS: invalid variable name
Error: Could not find or load main class 3.2.3
sylvia@sylvia-VirtualBox:~$ ls
279ex-2.txt  Downloads  Pictures  Templates  weatherreducer.py
data.json   hadoop-3.2.3  Public   Videos
Desktop     mapper.py    reducer.py weatherdata.csv
Documents  Music        snap     weathermapper.py
sylvia@sylvia-VirtualBox:~$ cat weatherdata.csv | python3 weathermapper.py | sor
t |python3 weatherreducer.py
Los Angeles    76.5
New York       31.0
sylvia@sylvia-VirtualBox:~$
```

Step 2: Mapper Logic - mapper.py:

Create a file named "mapper.py" to implement the logic for the mapper. The mapper will read input data from STDIN, split lines into words, and output each word with its count.

```

nano mapper.py
# Copy and paste the mapper.py code

#!/usr/bin/env python

import sys

# input comes from STDIN (standard input)
# the mapper will get daily max temperature and group it by month. so output will be
(month,daily_max_temperature)

for line in sys.stdin:
    # remove leading and trailing whitespace
    line = line.strip()
    # split the line into words
    words = line.split()
    # See the README hosted on the weather website which help us understand how each
    # position represents a column
    month = words[10]
    daily_max = words[11]
    # increase counters
    for word in words:
        # write the results to STDOUT (standard output);
        # what we output here will be go through the shuffle process and then
        # be the input for the Reduce step, i.e. the input for reducer.py
        #
        # tab-delimited; month and daily max temperature as output
        print('%s\t%s' % (month, daily_max))

```

Step 3: Reducer Logic - reducer.py:

Create a file named "reducer.py" to implement the logic for the reducer. The reducer will aggregate the occurrences of each word and generate the final output.

```

nano reducer.py
# Copy and paste the reducer.py code

```

reducer.py

```

#!/usr/bin/env python

from operator import itemgetter
import sys

#reducer will get the input from stdid which will be a collection of key, value(Key=month , value=
daily_max_temperature)
#reducer logic: will get all the daily max temperature for a month and find max temperature for the
month
#shuffle will ensure that key are sorted(month)

```

```

current_month = None
current_max = 0
month = None

# input comes from STDIN for
line in sys.stdin:
    # remove leading and trailing whitespace    line
    = line.strip()
    # parse the input we got from mapper.py    month,
    daily_max = line.split('\t', 1)

    # convert daily_max (currently a string) to float    try:
        daily_max = float(daily_max)    except
ValueError:
    # daily_max was not a number, so silently
    # ignore/discard this line
continue

# this IF-switch only works because Hadoop shuffle process sorts map output
# by key (here: month) before it is passed to the reducer
if current_month == month:    if daily_max > current_max:
    current_max = daily_max    else:    if current_month:
        # write result to STDOUT
        print ('%s\t%s' % (current_month, current_max))
    current_max = daily_max
    current_month = month

# output of the last month if current_month == month:
print ('%s\t%s' % (current_month, current_max))

```

Step 4: Prepare Hadoop Environment:

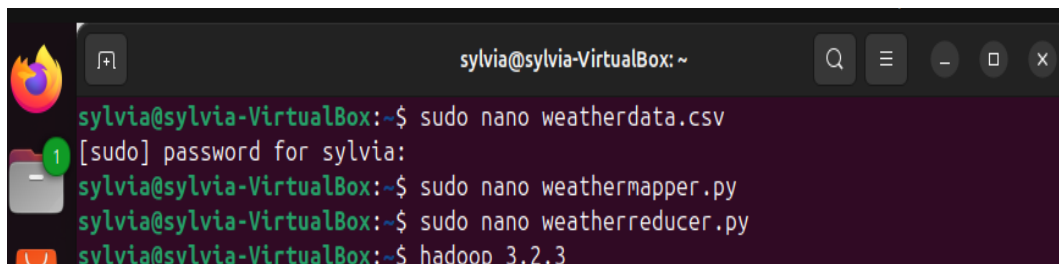
Start the Hadoop daemons and create a directory in HDFS to store your data.

```
start-all.sh
```

Step 6: Make Python Files Executable:

Give executable permissions to your mapper.py and reducer.py files.

```
chmod 777 mapper.py reducer.py
```



```
sylvia@sylvia-VirtualBox: ~
sylvia@sylvia-VirtualBox:~$ sudo nano weatherdata.csv
[sudo] password for sylvia:
sylvia@sylvia-VirtualBox:~$ sudo nano weathermapper.py
sylvia@sylvia-VirtualBox:~$ sudo nano weatherreducer.py
sylvia@sylvia-VirtualBox:~$ hadoop 3.2.3
```

Step 7: Run the program using Hadoop Streaming:

Download the latest hadoop-streaming jar file and place it in a location you can easily access.

Then run the program using Hadoop Streaming.

```
hadoop fs -mkdir -p /weatherdata
```

```
hadoop fs -copyFromLocal /home/sx/Downloads/dataset.txt /weatherdata
```

```
hdfs dfs -ls /weatherdata
```

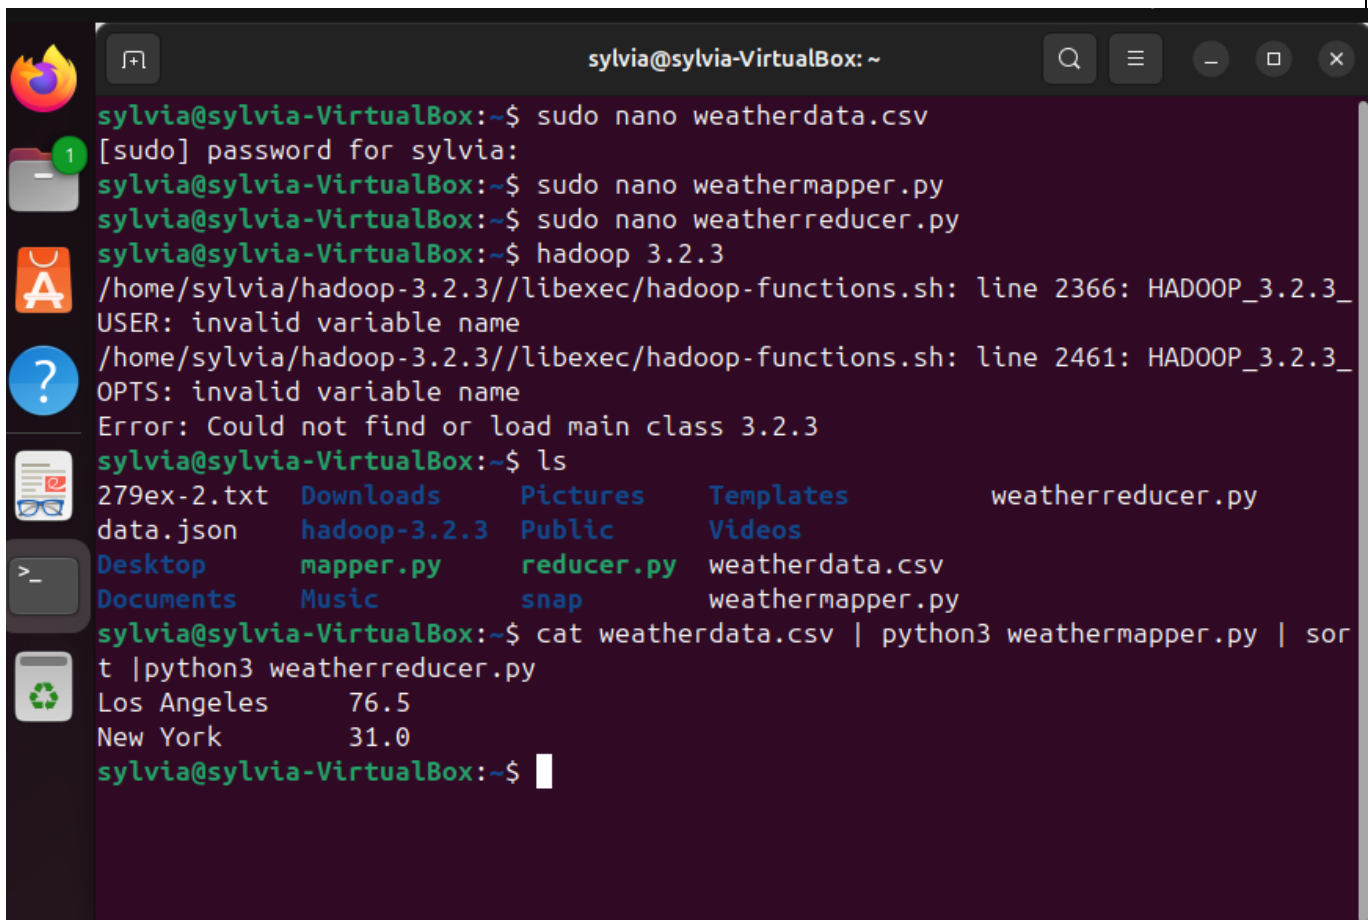
```
hadoop jar /home/sx/hadoop-3.2.3/share/hadoop/tools/lib/hadoop-streaming-3.2.3.jar \
-input /weatherdata/dataset.txt \
-output /weatherdata/output \
-file "/home/sx/Downloads/mapper.py" \
-mapper "python3 mapper.py" \
-file "/home/sx/Downloads/reducer.py" \
-reducer "python3 reducer.py"
```

```
hdfs dfs -text /weatherdata/output/* > /home/sx/Downloads/outputfile.txt
```

Step 8: Check Output:

Check the output of the program in the specified HDFS output directory.

```
hdfs dfs -text /weatherdata/output/* > /home/sx/Downloads/output/ /part-00000
```



```

sylvia@sylvia-VirtualBox: ~
sylvia@sylvia-VirtualBox:~$ sudo nano weatherdata.csv
[sudo] password for sylvia:
sylvia@sylvia-VirtualBox:~$ sudo nano weathermapper.py
sylvia@sylvia-VirtualBox:~$ sudo nano weatherreducer.py
sylvia@sylvia-VirtualBox:~$ hadoop 3.2.3
/home/sylvia/hadoop-3.2.3//libexec/hadoop-functions.sh: line 2366: HADOOP_3.2.3_
USER: invalid variable name
/home/sylvia/hadoop-3.2.3//libexec/hadoop-functions.sh: line 2461: HADOOP_3.2.3_
OPTS: invalid variable name
Error: Could not find or load main class 3.2.3
sylvia@sylvia-VirtualBox:~$ ls
279ex-2.txt  Downloads  Pictures  Templates  weatherreducer.py
data.json   hadoop-3.2.3  Public  Videos
Desktop     mapper.py    reducer.py  weatherdata.csv
Documents  Music        snap       weathermapper.py
sylvia@sylvia-VirtualBox:~$ cat weatherdata.csv | python3 weathermapper.py | sor
t |python3 weatherreducer.py
Los Angeles      76.5
New York         31.0
sylvia@sylvia-VirtualBox:~$

```

After copy and paste the above output in your local file give the below command to remove the directory from hdfs : `hadoop fs -rm -r /weatherdata/output`

Result:

Thus, the program for weather dataset using Map Reduce has been executed successfully.